

# Lecture 4: Optimization

CS 182 Spring 2021 – Taught by Sergey Levine

Notes by Michael Zhu

## Gradient Descent

General Algorithm:

1. Find a direction  $v$  where  $L(\theta)$  decreases
2. Update  $\theta \leftarrow \theta + \alpha v$

$$\text{Loss func: } L(\theta) = - \sum_i \log p_\theta(y_i|x_i)$$

$$\text{Optimization objective: } \theta^* \leftarrow \arg \min_{\theta} L(\theta)$$

For each dimension,  $L(\theta)$  decreases in the direction **opposite the slope**

We can compute the slope at each direction by taking the **partial derivative of the loss** with respect to each parameter  $dL(\theta)/d\theta_i$

This is also known as the **gradient**:

$$\nabla_{\theta} L(\theta) = \begin{pmatrix} dL(\theta)/d\theta_1 \\ dL(\theta)/d\theta_2 \\ \dots \\ dL(\theta)/d\theta_n \end{pmatrix}$$

## Issues with Gradient Descent

Gradient descent often works well for optimizing **convex** loss functions. Negative log likelihood loss for logistic regression is an example of a convex loss function.

However, the loss surface of a neural network is usually **non-convex**.

### Local optima

This is the most obvious issue with non-convex loss landscapes.

Gradient descent could converge to a solution that is **arbitrarily worse** than the global optimum.

Surprisingly, this local optima are not a huge issue for large networks. There is a lot of evidence that shows local optima in big neural nets are usually not that much worse than the global optima.<sup>1</sup>

### Plateaus

If your algorithm reaches a **plateau** with a **small learning rate**, it could get **stuck**.

**Momentum** helps us deal with plateaus.

---

<sup>1</sup>The Loss Surfaces of Multilayer Networks by Choromanska et al. <https://arxiv.org/pdf/1412.0233.pdf>

## Saddle Points

Definition: a point that is a **local minimum** in some dimensions and a **local maximum** in other dimensions

Gradients become very small (aka **vanish**) at saddle points; similar effect as plateaus that **slows** down gradient descent

**Very common for high dimensional data**

**Critical points:** where gradient is zero  $\nabla_{\theta} L(\theta) = 0$ ; can be a **maximum**, **minimum**, or **saddle**

For **2-dimensions**, saddle points don't exist; you can figure out if a critical point is a max or min using the **second derivative**

- **local maximum:**  $d^2 L / d\theta^2 < 0$
- **local minimum:**  $d^2 L / d\theta^2 > 0$

In **higher dimensions**, we can use the **Hessian matrix** to classify critical points

$$\begin{pmatrix} \frac{\partial^2 L}{\partial \theta_1 \partial \theta_1} & \cdots & \frac{\partial^2 L}{\partial \theta_1 \partial \theta_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 L}{\partial \theta_n \partial \theta_1} & \cdots & \frac{\partial^2 L}{\partial \theta_n \partial \theta_n} \end{pmatrix}$$

At a critical point, the **Hessian matrix** will be a **diagonal matrix**

- **maximum:** all diagonal entries are **positive**
- **minimum:** all diagonal entries are **negative**
- **saddle:** diagonal entries have **mismatched** signs

At higher dimensions, it's pretty **unlikely** for all the diagonal entries to have the same sign. This tells us that **saddle points** are the **most common** critical points at higher dimensions.

## Momentum

Gradient descent always moves in the direction of **steepest descent**

However, this is **NOT** always the most optimal direction (imagine a **zig-zag** in a valley shaped loss surface; its highly repetitive and slow)

Momentum is a **heuristic** that speeds up gradient descent by **adding past gradient steps**

$$\text{Update rule: } \theta_{k+1} = \theta_k - \alpha g_k$$

$$\text{Without momentum: } g_k = \nabla_{\theta} L(\theta_k)$$

$$\text{With momentum: } g_k = \nabla_{\theta} L(\theta_k) + \mu g_{k-1}$$

Possible scenarios (for intuition, not guaranteed to occur):

1. **Gradient steps zig-zag:** momentum causes “zigging” directions to **cancel** each other out and amplifies the magnitude of the **aggregate direction**
2. **Gradient steps DON'T zig-zag:** momentum causes gradient steps to **accelerate towards current direction** (useful for **plateaus**)

**Note:** momentum does not have many theoretical guarantees but it's a heuristic that generally works well in practice with virtually **no added computational cost**

**Note:** can look up “Nesterov accelerated gradient” which is a similar idea that **does** have some appealing guarantees in convex optimization, but for neural nets in practice we just use momentum

## RMSProp

**Issue:** gradients can have very extreme scaling (i.e. super small or super big) based on the difference between  $f_\theta(x)$  and  $y$

**Solution:** we can **normalize** the magnitude along each dimension

$$s_k \leftarrow \beta s_{k-1} + (1 - \beta)(\nabla_\theta L(\theta_k))^2$$
$$\theta_{k+1} = \theta_k - \alpha \frac{\nabla_\theta L(\theta_k)}{\sqrt{s_k}}$$

$s_k$  is essentially a **running average** of the magnitude of gradient at dimension  $k$

**Note:** we need to square  $\nabla_\theta L(\theta_k)$  to get magnitude so it's not affected by direction

Like **momentum**, RMSProp works pretty well in practice but doesn't have many theoretical guarantees

## AdaGrad

$$s_k \leftarrow s_{k-1} + (\nabla_\theta L(\theta_k))^2$$
$$\theta_{k+1} = \theta_k - \alpha \frac{\nabla_\theta L(\theta_k)}{\sqrt{s_k}}$$

Compared to RMSProp, AdaGrad is the **cumulative magnitude** per dimension instead of a running average

- Can be proven to work well for **convex** problems
- The learning rate  $\alpha/\sqrt{s_k}$  decreases over time, which is good for convex problems
- Only works for nonconvex problems if the learning algorithm can find the optimum region quickly before learning rate decays too much
  - RMSProp works better for **nonconvex** because it effectively “forgets” past gradients while AdaGrad never “forgets”

## Adam

**Idea:** use both momentum and RMSProp

$$m_k = (1 - \beta_1)\nabla_\theta L(\theta_k) + \beta_1 m_{k-1} \quad \text{first moment estimate (similar to "momentum")}$$

$$v_k = (1 - \beta_2)(\nabla_\theta L(\theta_k))^2 + \beta_2 v_{k-1} \quad \text{second moment estimate (similar to "RMSProp")}$$

$$\hat{m}_k = \frac{m_k}{1 - \beta_1^k}$$

$$\hat{v}_k = \frac{v_k}{1 - \beta_2^k}$$

$$\text{Update step: } \theta_{k+1} = \theta_k - \alpha \frac{\hat{m}_k}{\sqrt{\hat{v}_k} + \epsilon}$$

**Popular default hyperparameter values:**

- $\alpha = 0.001$
- $\beta_1 = 0.9$
- $\beta_2 = 0.999$

$m_k$  and  $v_k$  start at 0 so will be numerically small for the first few iterations

$\hat{m}_k$  and  $\hat{v}_k$  are used to speed up the first few iterations by dividing by  $1 - \beta^k$

# Stochastic Gradient Descent

Regular gradient descent is very expensive for large datasets

$$L(\theta) = -\frac{1}{N} \sum_i^N \log p_{\theta}(y_i|x_i)$$

Each gradient step requires summing over **all datapoints** in the dataset

**Idea:** break the dataset into separate **batches** of size **B** which we use for each gradient step. This works if each batch is **randomly sampled** because they are an **unbiased approximation** of the expected loss

$$L(\theta) = -\frac{1}{N} \sum_i^N \log p_{\theta}(y_i|x_i) \approx -E_{p(x_i, y_i)}[\log p_{\theta}(y_i|x_i)] \approx -\frac{1}{B} \sum_j^B \log p_{\theta}(y_{i_j}|x_{i_j})$$

## SGD with minibatches procedure

1. Sample  $B \subset D$
2. Estimate  $g_k \leftarrow -\nabla_{\theta} \sum_i^B \log p_{\theta}(y_i|x_i) \approx \nabla_{\theta} L(\theta)$
3. Update  $\theta_{k+1} \leftarrow \theta_k - \alpha g_k$ ; apply momentum, ADAM, RMSProp, etc.

Each iteration uses a different **minibatch**

Every pass through the entire dataset is an **epoch**

**For efficient memory access:**

- In practice, we randomly shuffle the dataset once at the start and then index into each minibatch sequentially (recall *principle of locality*)
- Batch sizes  $B$  are generally powers of two  $\rightarrow 1, 2, 4, 8, 16, 32, 64, 128, \dots$

## Tuning SGD

### Hyperparameters

- batch size  $B$ : larger batches  $\rightarrow$  less noise in gradients but more memory expensive
- learning rate  $\alpha$ : ideally use largest learning rate that doesn't diverge; decay over time if necessary
- momentum  $\mu$ : 0.99 is good default
- Adam  $\beta_1, \beta_2$ : 0.9 and 0.999 are good defaults usually

Technically, we want to tune these **optimization hyperparameters** on the **training loss**

In practice, people often tune on **validation loss** as well (recall validation loss is usually used for **regularization hyperparameters**)

Stochastic gradient descent can sometimes have a regularizing effect since each batch keeps the other batches from overfitting

### Learning rate decay

- Learning rate too high  $\rightarrow$  learning algorithm either diverges or keeps jumping around the optimum
- Learning rate too low  $\rightarrow$  increases training time and can get stuck on plateaus

**Idea: Start with a high learning rate** that doesn't diverge (avoid plateaus and quickly get near optimum); **gradually decrease** it to get closer to optimum

Learning rate decay is usually needed for best performance from SGD+momentum

Often not needed with ADAM